


# Project III | AI ChatBot for YouTube Video QA and summary

Nataliya Demyanenko



# Table of contents

01

**Project Scope**

---

02

**App**

---

03

**Other Steps**

---

04

**Conclusion**

---



**01**

# **Project Scope**

---

# Project Scope

---

Goal: Build a chatbot for YouTube video summarization and Q&A

Focus areas:

- Exploration of open-source solutions
- Optimization for CPU deployment
- Abstractive question answering
- Real-life processing of user-provided videos

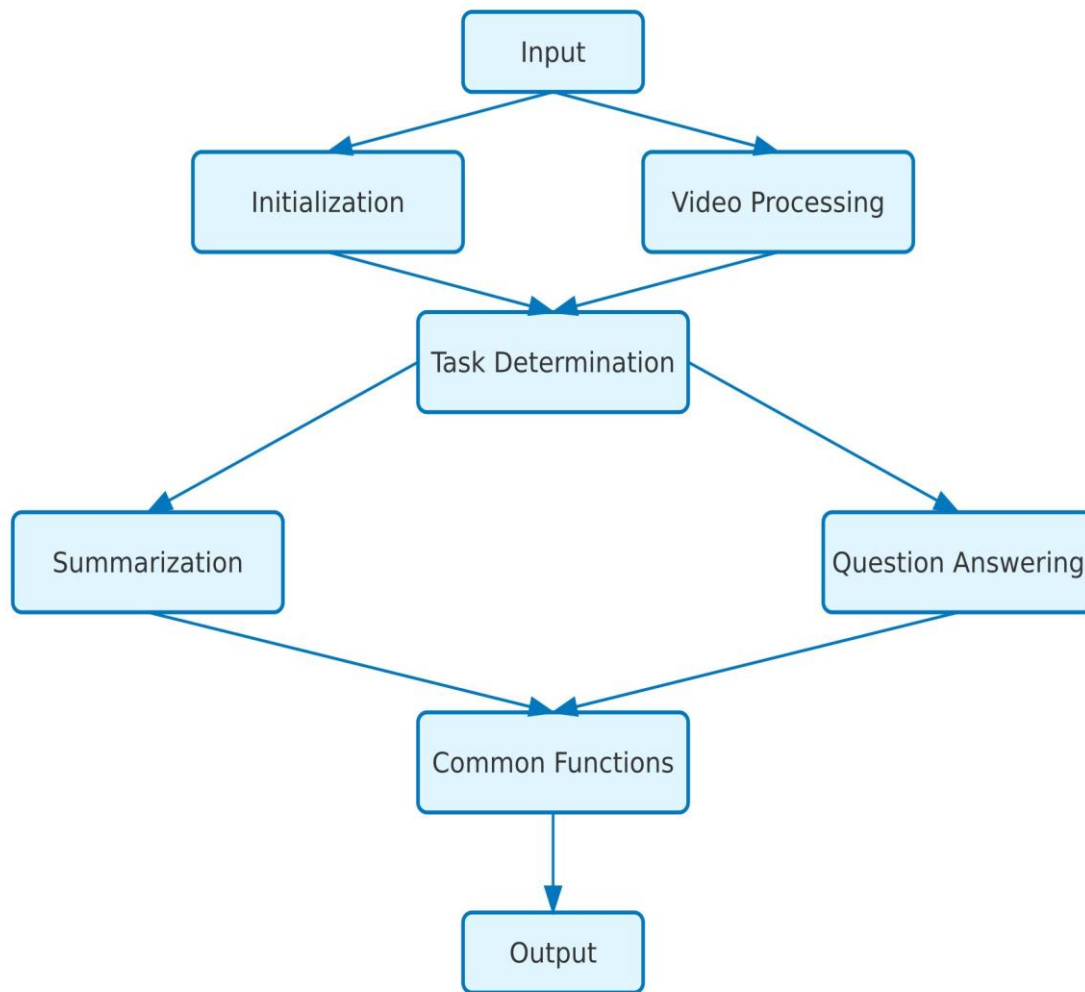




**02**

**App**

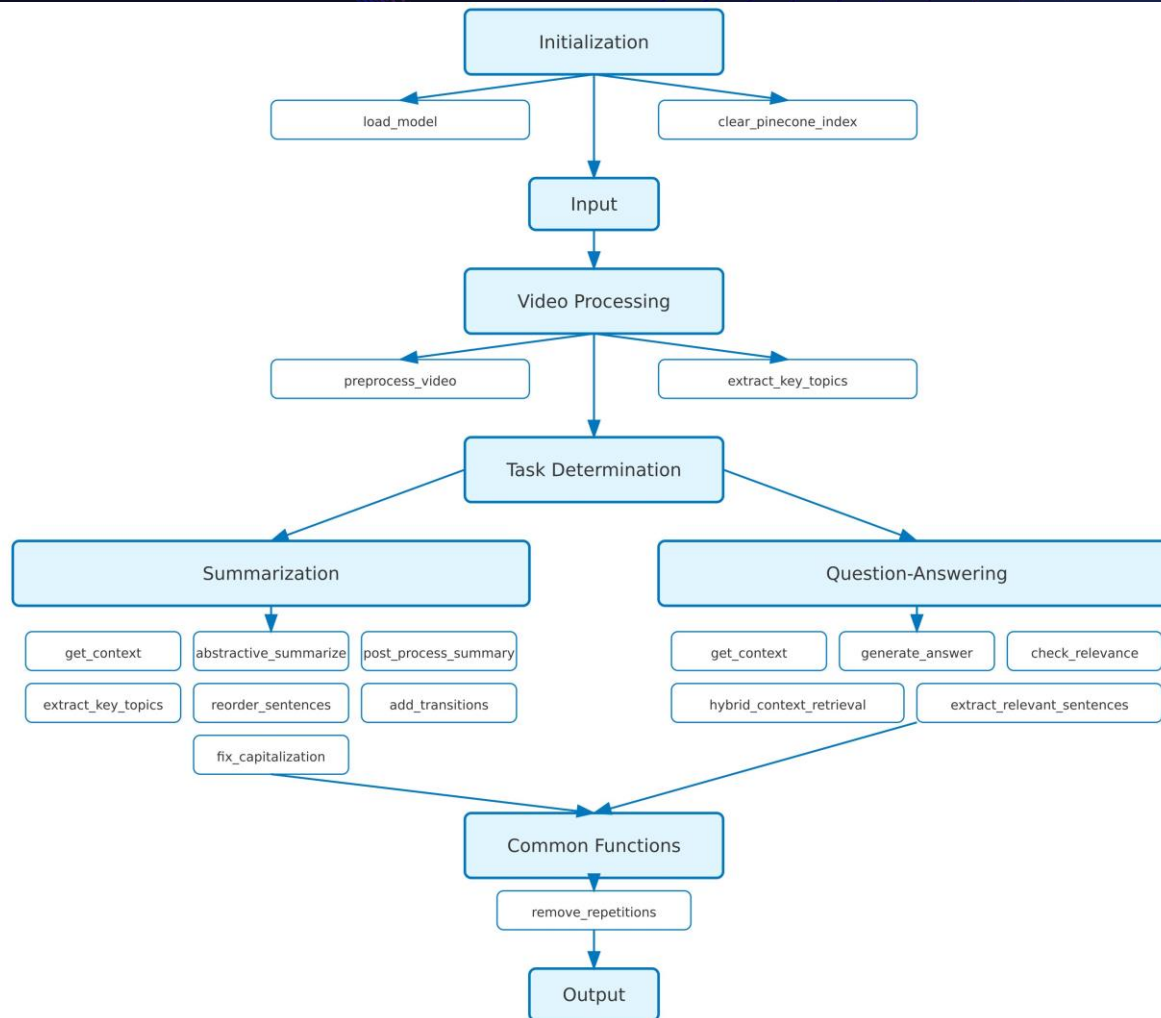
---



# General overview of the app

1. **Initialization:** model loading and index clearing.
2. **Video Processing:** Extraction and chunking of video content.
3. **Task Determination:** Q&A or summary
4. **Summarization:** Multi-step process involving context retrieval, abstractive summarization, and post-processing.
5. **Q&A:** Utilizes hybrid context retrieval and relevance checking for accurate answers.
6. **Common Functions:** Post-processing for output quality.

# All functions



# Video Preprocessing



**URL upload**



- Extracts video information and transcript
- Splits the content into manageable chunks
- Upserts to Pinecone

```
video_data.append({  
    'chunk_id': chunk_id,  
    'video_id': video_id,  
    'title': info['title'],  
    'channel': info['uploader'],  
    'description': info['description'],  
    'duration': info['duration'],  
    'start_time': combined_start_time,  
    'end_time': combined_start_time + combined_duration,  
    'source': f"https://www.youtube.com/watch?v={video_id}&t={int(combined_start_time)}",  
    'chunk_text': text,  
    'metadata_context': f"Video Title: {info['title']}\n"  
})
```



# Task Determination



## Function: `determine_task()`

Purpose: Decide between summarization and Q&A

Process:

- Checks user input against predefined summarization phrases
- Classifies as 'summarization' or 'question\_answering'

Advantage: Easily expandable for future task types

# List of phrases that indicate a summary request

`summary_phrases =`


```
["summarize", "summary", "summarize the video", "what  
is this video about", "what is the video about", "what is  
the main topic of the video", "what's the video about",  
"what's the video about", "give me an overview", "brief  
overview", "overview", "main points", "key points",  
"main idea", "what is the video discussing"]
```

# Model Selection for Q&A

**Selected model: VBlagoje/BART-LFQA**

Reasons for selection:

- Advanced question-answering capabilities
- State-of-the-art language model (BART architecture)
- Pre-trained on large datasets
- Customization and fine-tuning potential
- Optimized for long-form answers



# Abstractive Summarization

Model Used: facebook/bart-large-cnn


Function: `abstractive_summarize()`

Process:

- Prepend key topics to guide summary
- Use beam search, length penalties, and sampling

Parameters:

- Beam size, max/min length, length penalty, etc.



# Hybrid Context Retrieval

## Function: `hybrid_context_retrieval()`

Combines TF-IDF and semantic similarity for more accurate context retrieval.

- Compute TF-IDF scores
- Generate semantic embeddings
- Combine scores with weighted average

Advantage: Balances keyword matching and semantic understanding

```
def hybrid_context_retrieval(query, contexts, top_k=7, tfidf_weight=0.4):  
  
    # TF-IDF  
    vectorizer = TfidfVectorizer(stop_words='english')  
    tfidf_matrix = vectorizer.fit_transform([query] + contexts)  
    tfidf_scores = cosine_similarity(tfidf_matrix[0:1], tfidf_matrix[1:]).flatten()  
  
    # Semantic Similarity  
    query_embedding = retriever.encode([query])  
    context_embeddings = retriever.encode(contexts)  
    semantic_scores = cosine_similarity(query_embedding, context_embeddings)[0]  
  
    # Combine scores  
    combined_scores = 0.5 * tfidf_scores + 0.5 * semantic_scores  
    top_indices = combined_scores.argsort()[-top_k:][::-1]  
    return [contexts[i] for i in top_indices]
```

# Post-Processing and Formatting



## Key Functions:

- `post_process_summary()`
  - `remove_repetitions()`
  - `add_transitions()`
  - `fix_capitalization()`
- 

Purpose: Enhance readability and coherence

Techniques:

- Sentence reordering
- Redundancy removal
- Proper capitalization
- Transition phrase insertion



**03**

# Other Steps

---

# Memory implementation



**Type:**  
**ConversationBuffer**  
**WindowMemory**

---

Purpose: Maintain context of recent interactions  
Implementation:

- Set window size ( $k=3$ )
- Integrate with ConversationChain
- Implemented in Jupyter Notebook, not in deployed app

Advantage: Enables more coherent multi-turn conversations



# Fine-tuning with PEFT-optimization & Quantization

Due to time constraints, computational limits, and the size of the model being fine-tuned, I had to drastically limit training parameters, which resulted in a poor performance:

	vblagoje/bart_lfqa	Model fine-tuned on SQUAD
<b>Faithfulness</b>	0.653	0.484
<b>Context Recall</b>	0.578	0.578
<b>Relevance</b>	0.110	0.110
<b>ROUGE-1</b>	0.110	0.110
<b>ROUGE-2</b>	0.056	0.056
<b>ROUGE-L</b>	0.110	0.110

I attempted to fine tune the model on two datasets, the Stanford Question Answering Dataset (SQuAD) and ELI5-Category dataset, that both seemed promising fit.





**04**

# **Conclusion**

---

# Conclusion and Future work



## Key Takeaways:

---

More cautious approach to project scope definition



## Future Improvements

---

- Processing videos without YouTube transcript
- Fine-tuning both Q&A and summarization models
- Advanced sentence reordering in summarization
- Integration of sophisticated NLP techniques



# Thanks!

**CREDITS:** This presentation template was created by **Slidesgo**, including icons by **Flaticon** and infographics & images by **Freepik**

Please keep this slide for attribution